



Automatic Parallelization & Manual Parallelization with Open MP

Shrikant Vinchurkar
Mayank Chaudhary

Results from blocking

- Horizontal blocking into 4 blocks
- Interleaving of iterations
- Removed copy operation in relax_jacobi
- Time reduced from 11 sec to 8.76
- We are continuing OpenMP parallelization with this code

Interleaving approach

1--> 3

1--> 3

1--> 3

1--> 3

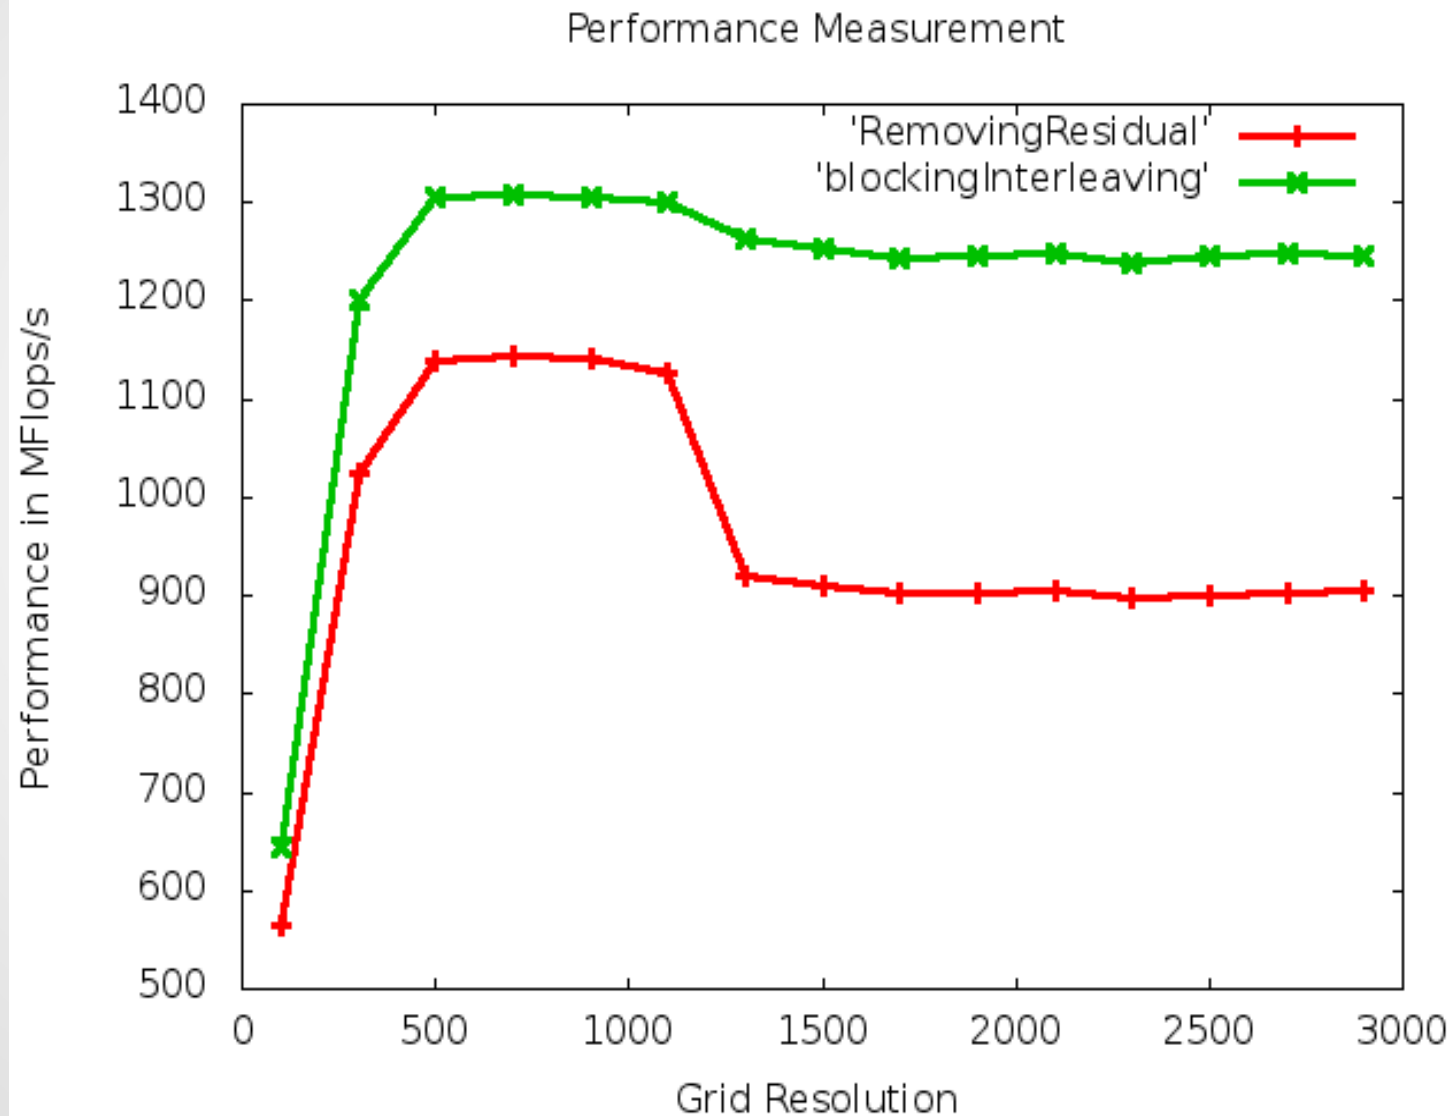
0->2->4

0->2->4

0->2->4

0->2->4

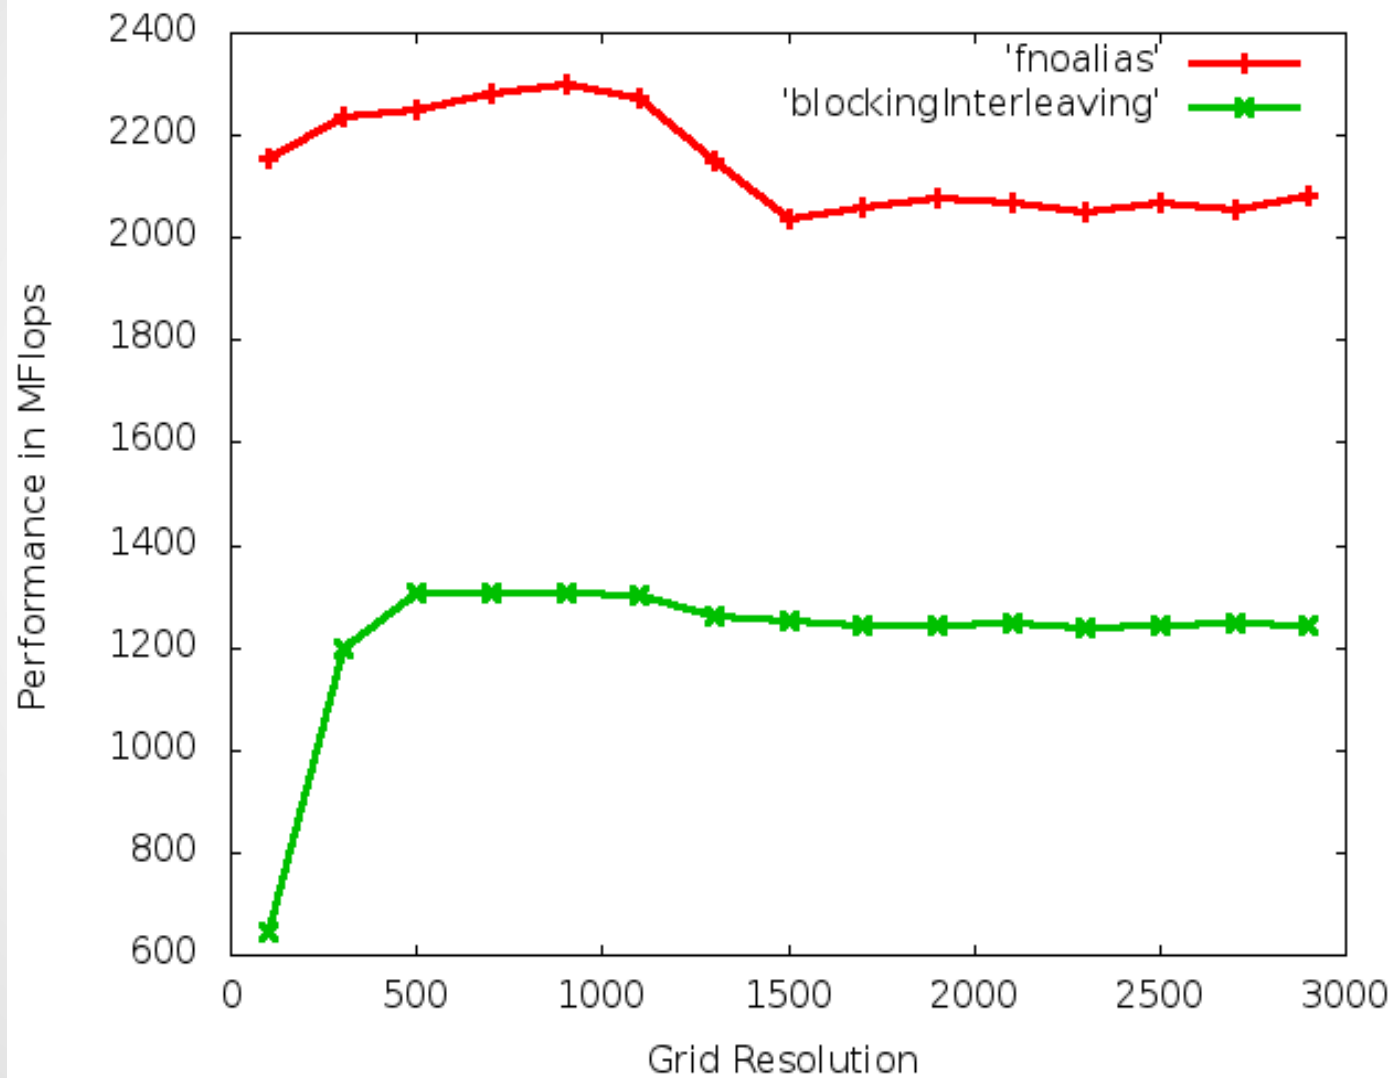
Performance Improvement



No Aliasing boost

- `-fargument-noalias` : option tell compiler that function arguments cannot alias each other
- It gave a huge boost, runtime down to 6 sec from 9 sec

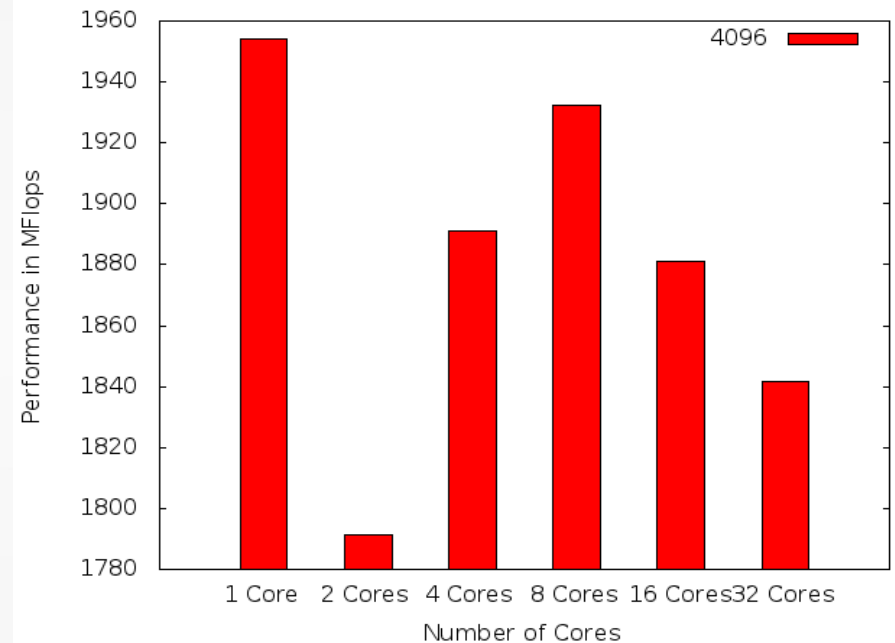
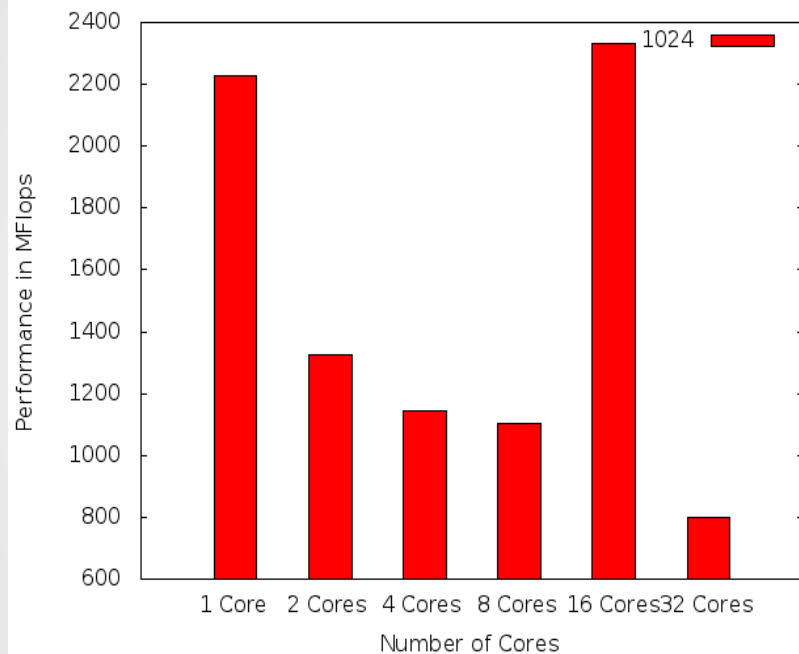
Fnoalias boost



Automatic parallelization pragmas

- `#pragma parallel`
- `#pragma loop count min(*)`
- `#pragma ivdep`
- `#pragma parallel always`

Automatic parallelization results




Why Automatic parallelization failed?

- Used `-par-report3` to generate autoparallelization report
- Vector dependencies, tried to remove with `#pragma ivdep`
- “Loop not parallelized:insufficient computational work”
- Removed with “`-par-threshold[n]`”, $n=\{0..100\}$
- “Loop not parallelized:insufficient inner loop”
- Removed with `#pragma parallel always`
- These workarounds resulted in overhead with autoparallelization =>degrades performance


OpenMP Manual Parallelization

- `omp_set_num_threads(num)`
- `#pragma omp parallel`
- `#pragma omp for private(diff) reduction(+:sum)`

OpenMP Boost



OMP performance on different #cores



NUMA first touch allocation

- Modern operating systems all use virtual memory
- The OS typically optimizes memory allocations
 - malloc() does not allocate the memory directly
 - Only the memory management “knows” about the memory allocation, but no memory pages are made available
 - At the first memory access, the OS physically allocates the corresponding page (First Touch Policy)
- On NUMA systems this might lead to performance issues in threaded or multi-process applications...

(Reference: Numa Optimization(Intel), bit.ly/13okCjW)

Changes in code

```
// initialize data  
for (size_t i = 0; i < N; i++)  
for (size_t j = 0; j < M; j++) {...}
```

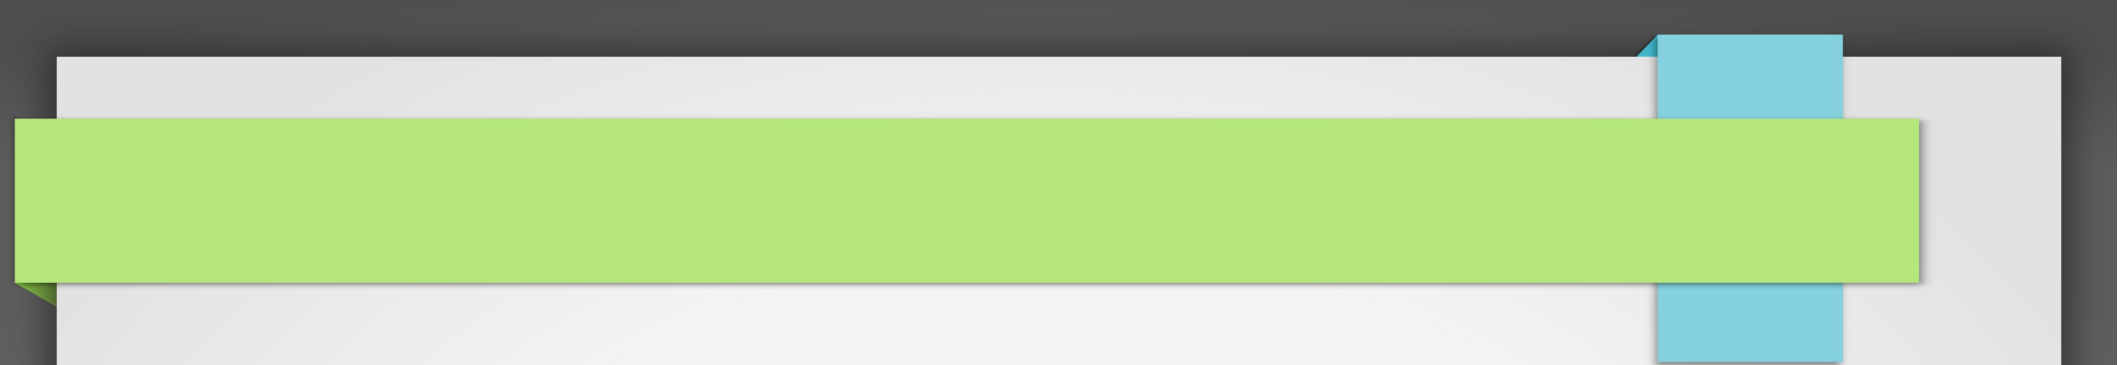
```
// perform work  
#pragma omp parallel for private(j)  
for (size_t i = 0; i < N; i++)  
for (size_t j = 0; j < M; j++) {...}
```

CHANGED TO

```
// initialize data  
#pragma omp parallel for private(j)  
for (size_t i = 0; i < N; i++)  
for (size_t j = 0; j < M; j++) {...}  
// perform work  
#pragma omp parallel for private(j)  
for (size_t i = 0; i < N; i++)  
for (size_t j = 0; j < M; j++) {...}
```

**Peak performance for 32
Threads increased from
7237 to 9563**

(Reference: Numa Optimization(Intel), bit.ly/13okCjW)

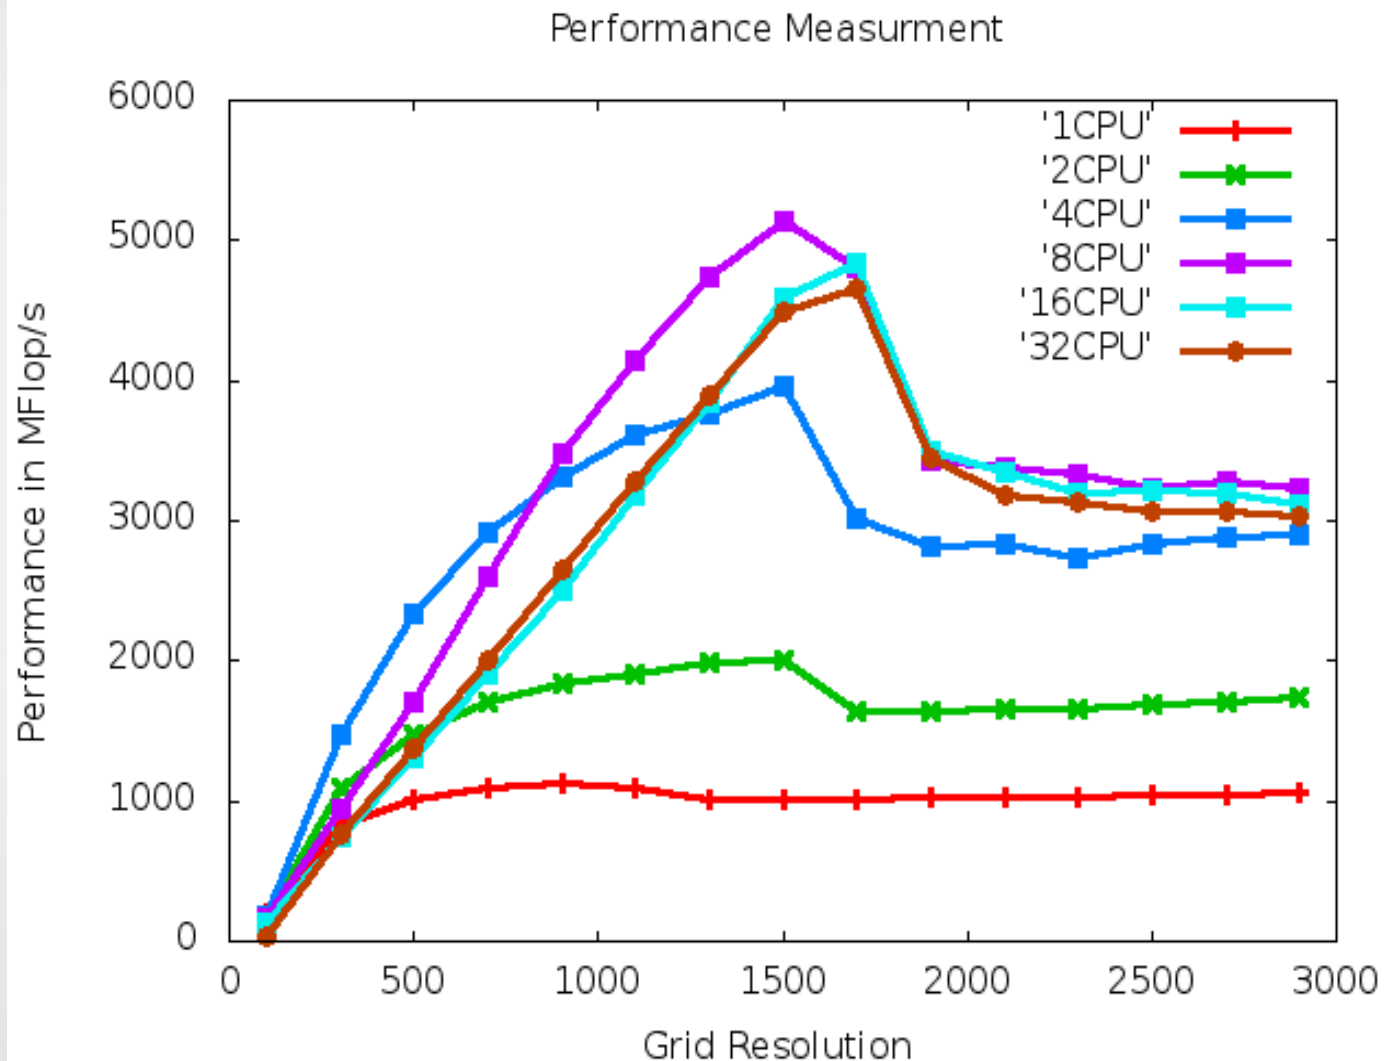


New Readings of Auto Parallelization, NonNuma & Numa


Auto Parallelization

- Auto Parallelization done on code which differs from original code with respect to following:
 - Loop interchange
 - Avoid copy operation
 - Change in the way residual is calculated


Auto Parallelization with grid resolution 100-2900



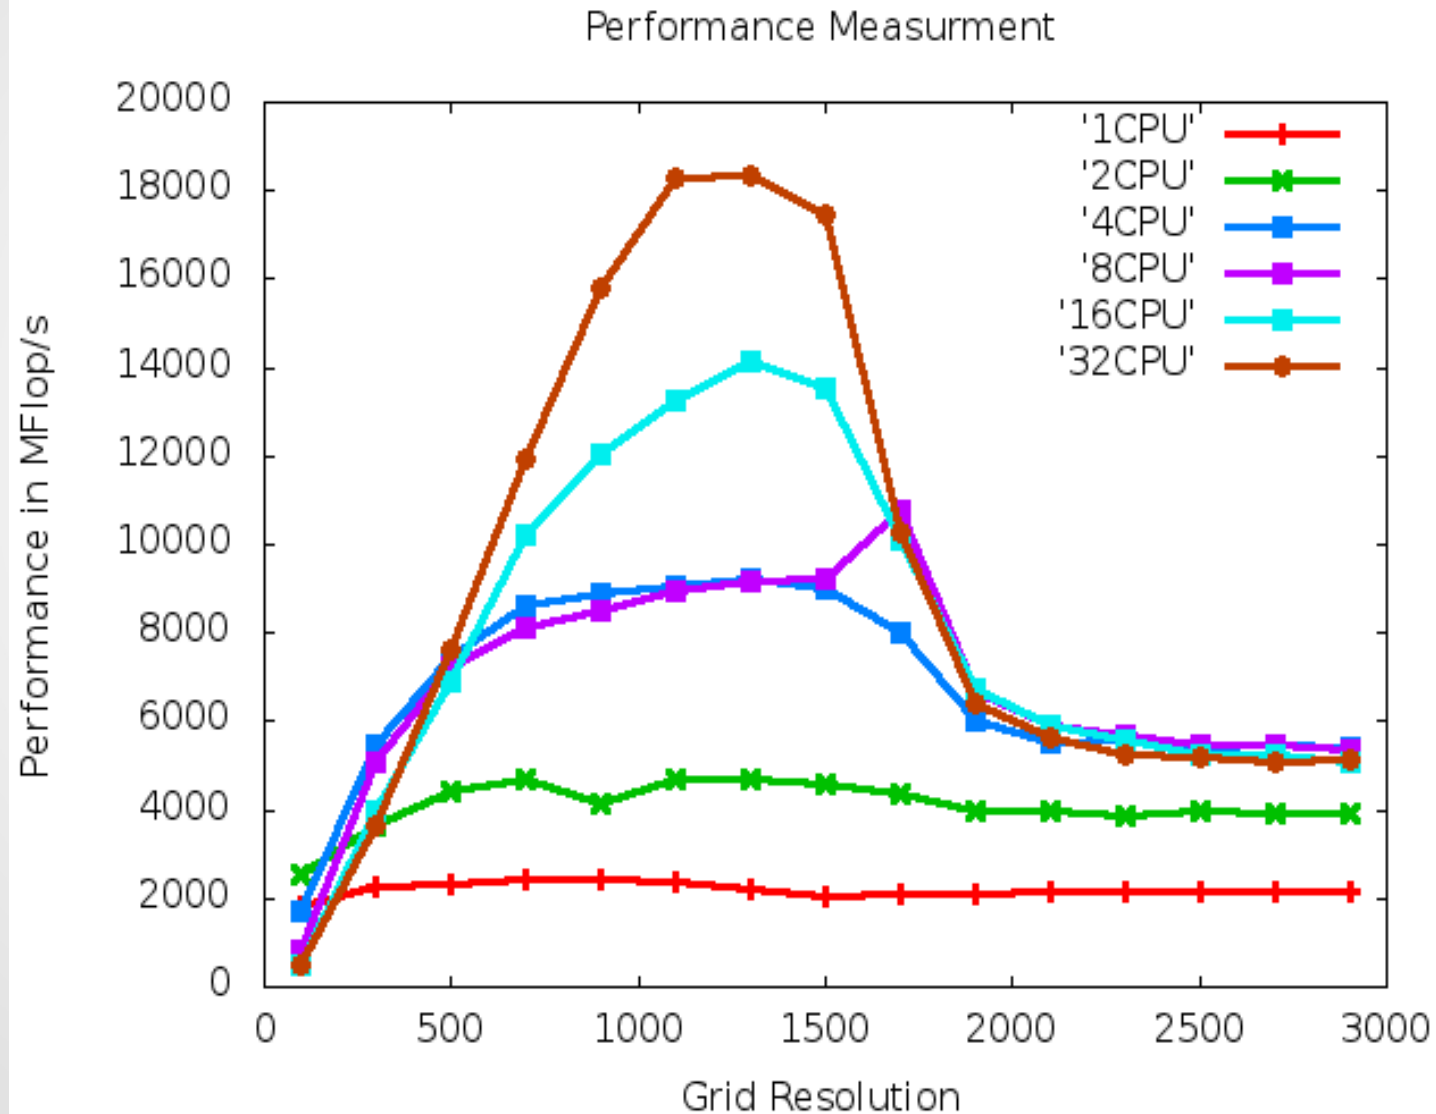
Auto Parallelize for grid 1024 & 4096




OpenMp performance without NUMA



OpenMP with NUMA



OpenMP with given grid points, non Numa



OpenMP with given grid points, with Numa

